

Annotated  
Version

Machine Learning Course - CS-433

# Support Vector Machines

Oct 29, 2019

changes by Martin Jaggi 2019, changes by Rüdiger Urbanke 2018, changes by Martin Jaggi  
2016, 2017 ©Mohammad Emtiyaz Khan 2015

Last updated on: October 29, 2019



# Motivation

By changing the cost function of a linear classifier from Logistic to Hinge, we obtain the support vector machine (SVM).

## Support Vector Machine

Throughout, we will work with a classification problem and assume<sup>a</sup> that the labels  $y_n \in \{\pm 1\}$ .

(This is in contrast to logistic regression, where we have used the convention  $y_n \in \{0, 1\}$ .)

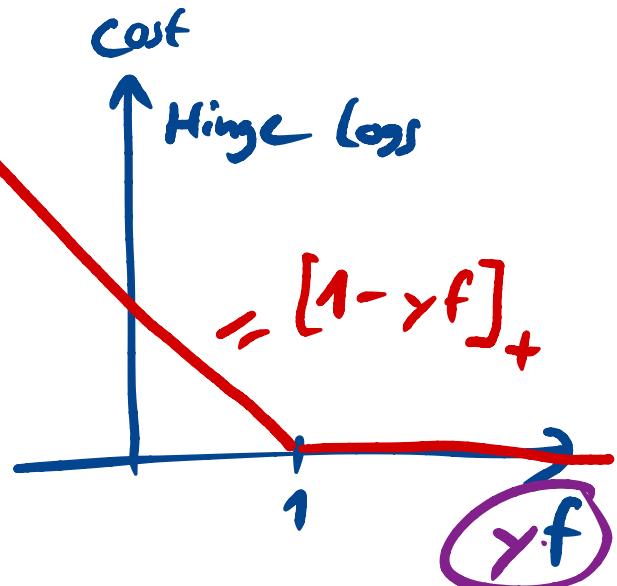
We again write  $\mathbf{x}_n$  for datapoint  $n$ , and assume that all constructed features and a potential constant bias are already included in  $\mathbf{x}_n$ .

The SVM optimizes the following cost:

$$\min_{\mathbf{w}} \sum_{n=1}^N [1 - y_n \mathbf{x}_n^\top \mathbf{w}]_+ + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

*label  $y_n \in \{+1, -1\}$*

where the first term is the Hinge loss defined as  $[z]_+ := \max\{0, z\}$ .



$$f = \mathbf{x}_n^\top \mathbf{w}$$

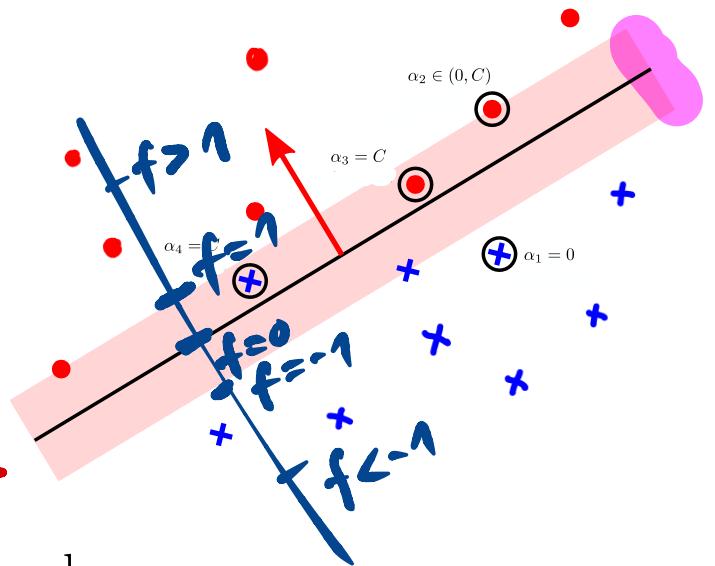
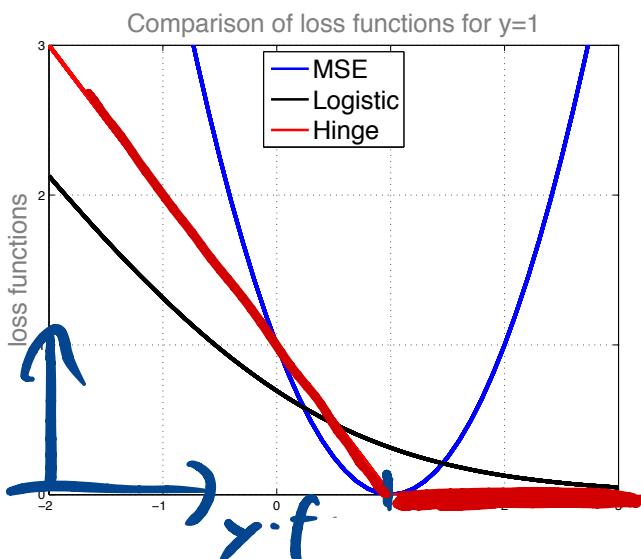
<sup>a</sup>Note that for any use-case, the labels can be converted accordingly before training, and after prediction.

# Hinge vs MSE vs Logistic

Consider  $y \in \{-1, +1\}$  with prediction  $f \in \mathbb{R}$ , then the three cost functions can be written as follows:

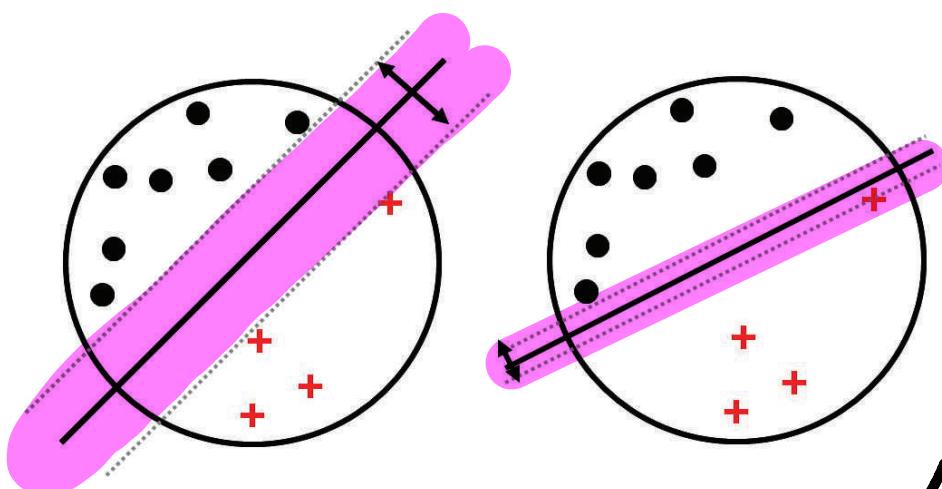
$$\begin{aligned} \text{Hinge}(f) &= [1 - yf]_+ \\ \text{MSE}(f) &= (1 - yf)^2 \\ \text{logisticLoss}(f) &= \log(1 + e^{-yf}) \end{aligned}$$

**SVM**  
**Regression for Classification**  
**Logistic Regression**  
**Homework**  $y$  transform



Notice the **margin** in the Hinge loss.  
SVM is a **maximum margin** method.

**Assumption :**  
**linearly separable**

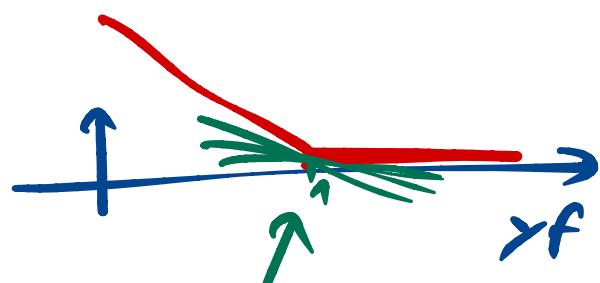


Margin  
 $\approx \frac{1}{\|w\|}$   
(not proven)

# Optimization

in  $\mathbf{w}$ : yes

Is this function convex? Is it differentiable? (in  $\mathbf{w}$ ) no!



$$\min_{\mathbf{w}} \sum_{n=1}^N [1 - y_n \mathbf{x}_n^\top \mathbf{w}]_+ + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

Can use SGD! (with subgradients).

Is there a better optimization algorithm here?

## Duality: The big picture

Let us say that we are interested in optimizing a function  $\mathcal{L}(\mathbf{w})$  and it is a difficult problem. Define an auxiliary function  $G(\mathbf{w}, \boldsymbol{\alpha})$  such that

$$\mathcal{L}(\mathbf{w}) = \max_{\boldsymbol{\alpha}} G(\mathbf{w}, \boldsymbol{\alpha}).$$

so that we can then choose between optimizing either of

$$\min_{\mathbf{w}} \max_{\boldsymbol{\alpha}} G(\mathbf{w}, \boldsymbol{\alpha}) = \max_{\boldsymbol{\alpha}} \min_{\mathbf{w}} G(\mathbf{w}, \boldsymbol{\alpha})$$

Primal problem      Dual problem

Three questions:

1. How do you set  $G(\mathbf{w}, \boldsymbol{\alpha})$ ?
2. When is it OK to switch  
 $\min_{\mathbf{w}}$  and  $\max_{\boldsymbol{\alpha}}$ ?
3. When is the dual easier to optimize than the primal?

**Q1:** How to obtain  $G(\mathbf{w}, \boldsymbol{\alpha})$ ?

For one datapoint

$$\begin{aligned} [v_n]_+ &:= \max\{0, v_n\} \\ &\stackrel{\Delta}{=} \max_{\alpha_n} \alpha_n v_n \text{ where } \alpha_n \in [0, 1] \\ [1 - y_n \mathbf{x}_n^\top \mathbf{w}]_+ &= \max_{\alpha_n \in [0, 1]} \underbrace{\alpha_n(1 - y_n \mathbf{x}_n^\top \mathbf{w})}_{v_n} \end{aligned}$$

*hinge-loss*

proof:

- case  $v \geq 0$   
 $\max = v$   
attained when  $\alpha = 1$
- case  $v < 0$   
 $\max = 0$   
attained at  $\alpha = 0$

For all points:

We can rewrite the SVM problem as:

$$\min_{\mathbf{w}} \max_{\boldsymbol{\alpha} \in [0, 1]^N} \sum_{n=1}^N \underbrace{\alpha_n(1 - y_n \mathbf{x}_n^\top \mathbf{w})}_{=: G(\mathbf{w}, \boldsymbol{\alpha})} + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

*L(w)*  
 $= \max_{\boldsymbol{\alpha}} G(\mathbf{w}, \boldsymbol{\alpha})$

This is differentiable, convex in  $\mathbf{w}$   
and concave in  $\boldsymbol{\alpha}$ .

over both  $\mathbf{w}$  and  $\boldsymbol{\alpha}$

**Q2:** When is it OK to switch max and min? Using a **minimax theorem**, it is OK to do so when  $G(\mathbf{w}, \boldsymbol{\alpha})$  is convex in  $\mathbf{w}$  and concave in  $\boldsymbol{\alpha}$ , under weak additional assumptions.

always  $\Leftrightarrow$

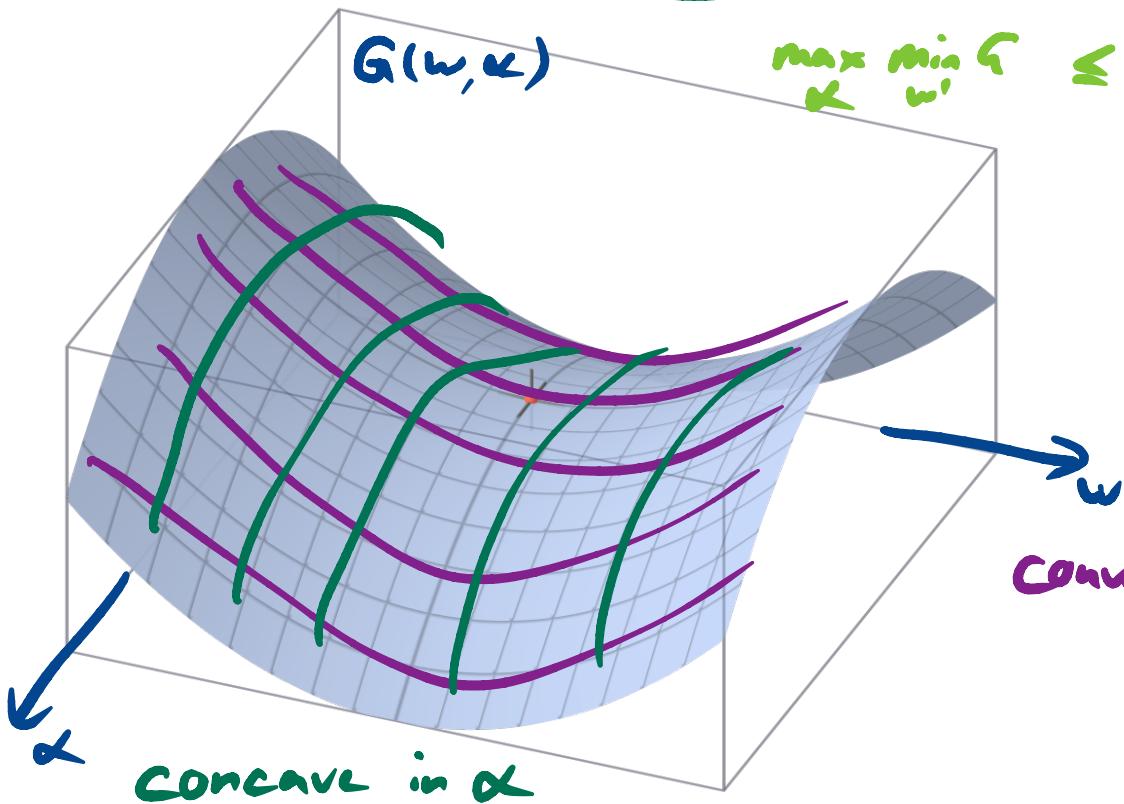
Proof

$$\max_{\boldsymbol{\alpha}} \min_{\mathbf{w}} G(\mathbf{w}, \boldsymbol{\alpha}) = \min_{\mathbf{w}} \max_{\boldsymbol{\alpha}} G(\mathbf{w}, \boldsymbol{\alpha})$$

$$\min_{\mathbf{w}'} G(\mathbf{w}', \boldsymbol{\alpha}) \leq G(\mathbf{w}, \boldsymbol{\alpha})$$

$$\max_{\boldsymbol{\alpha}} \min_{\mathbf{w}'} G(\mathbf{w}', \boldsymbol{\alpha}) \leq \max_{\boldsymbol{\alpha}} G(\mathbf{w}, \boldsymbol{\alpha})$$

$$\max_{\boldsymbol{\alpha}} \min_{\mathbf{w}'} G \leq \min_{\mathbf{w}} \max_{\boldsymbol{\alpha}} G(\mathbf{w}, \boldsymbol{\alpha})$$



For a more systematic way to derive suitable  $G(\mathbf{w}, \boldsymbol{\alpha})$  and dual variables  $\boldsymbol{\alpha}$ , see the concept of **convex conjugate** functions, as in the language of **Fenchel duality**.

See e.g. Bertsekas' "Nonlinear Programming" for more formal details.

For SVM, switching the min and max, we have the following saddle-point formulation

$$\max_{\alpha \in [0,1]^N} \min_{\mathbf{w}} \sum_{n=1}^N \alpha_n (1 - y_n \mathbf{x}_n^\top \mathbf{w}) + \frac{\lambda}{2} \|\mathbf{w}\|^2 \quad (1)$$

fixed

Taking the derivative w.r.t.  $\mathbf{w}$ :

$$\nabla_{\mathbf{w}} G(\mathbf{w}, \alpha) = - \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n + \lambda \mathbf{w} \stackrel{!}{=} 0$$

Equating this to  $\mathbf{0}$  (which is called the first-order optimality condition for  $\mathbf{w}$ ), we have the correspondence

$$\mathbf{w}(\alpha) = \frac{1}{\lambda} \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n = \frac{1}{\lambda} \mathbf{X}^\top \mathbf{Y} \alpha$$

where  $\mathbf{Y} := \text{diag}(\mathbf{y})$ , and  $\mathbf{X}$  again collects all  $N$  data examples as its rows.

Plugging this  $\mathbf{w} = \mathbf{w}(\alpha)$  back into the saddle-point formulation (1), we have the **dual optimization problem**:

$$\begin{aligned} & \max_{\alpha \in [0,1]^N} \sum_{n=1}^N \alpha_n \left( 1 - \frac{1}{\lambda} y_n \mathbf{x}_n^\top \mathbf{X}^\top \mathbf{Y} \alpha \right) + \frac{\lambda}{2} \left\| \frac{1}{\lambda} \mathbf{X}^\top \mathbf{Y} \alpha \right\|^2 \\ &= \max_{\alpha \in [0,1]^N} \alpha^\top \mathbf{1} - \frac{1}{2\lambda} \alpha^\top \underbrace{\mathbf{Y} \mathbf{X} \mathbf{X}^\top \mathbf{Y}}_{\text{matrix}} \alpha \end{aligned}$$

Dual problem  
of sum

**Q3:** When is the dual easier to optimize than the primal, and why?

- (1) The dual is a differentiable (but constrained) quadratic problem.

$$\max_{\boldsymbol{\alpha} \in [0,1]^N} \boldsymbol{\alpha}^\top \mathbf{1} - \frac{1}{2\lambda} \boldsymbol{\alpha}^\top \mathbf{Q} \boldsymbol{\alpha},$$

where  $\mathbf{Q} := \text{diag}(\mathbf{y}) \mathbf{X} \mathbf{X}^\top \text{diag}(\mathbf{y})$ .

Optimization is easy by using **co-ordinate descent**, or more precisely coordinate ascent since this is a maximization problem. Crucially, this method will be changing only one  $\alpha_n$  variable a time.

-  (2) The dual is naturally kernelized (just like the kernelized ridge, see next lecture) with  $\mathbf{K} := \mathbf{X} \mathbf{X}^\top$ .

- (3) The solution  $\boldsymbol{\alpha}$  is typically sparse, and is non-zero only for the training examples that are instrumental in determining the decision boundary.

Recall that  $\alpha_n$  is the slope of lines that are lower bounds to the Hinge loss.

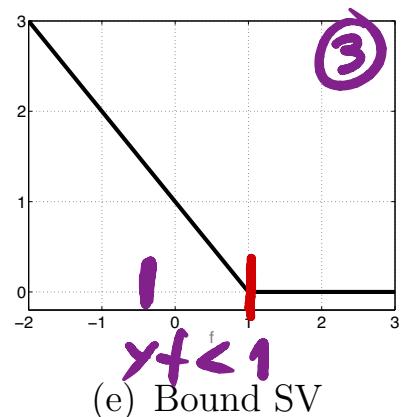
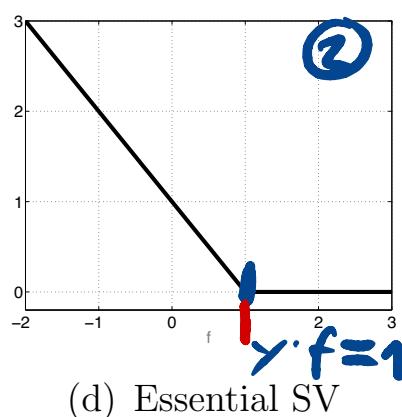
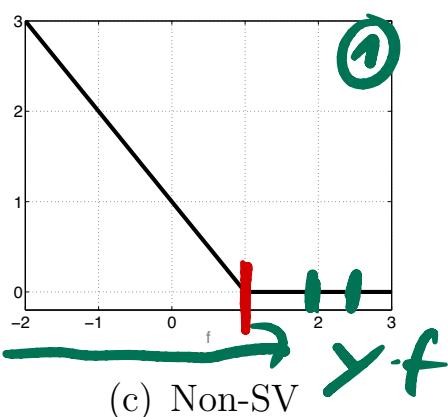
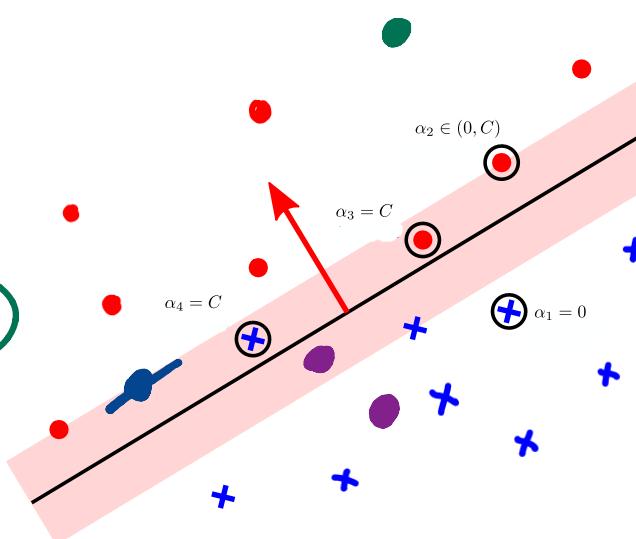
$$[1 - y_n f_n]_+ = \max_{\alpha_n \in [0, 1]} \alpha_n (1 - y_n f_n)$$

There are 3 kinds of data vectors  $\mathbf{x}_n$ .

1. Non-support vectors. Examples that lie on the correct side outside the margin, so  $\alpha_n = 0$ .

2. Essential support vectors. Examples that lie just on the margin, therefore  $\alpha_n \in (0, 1)$

3. Bound support vectors. Examples that lie strictly inside the margin, or on the wrong side, therefore  $\alpha_n = 1$ .

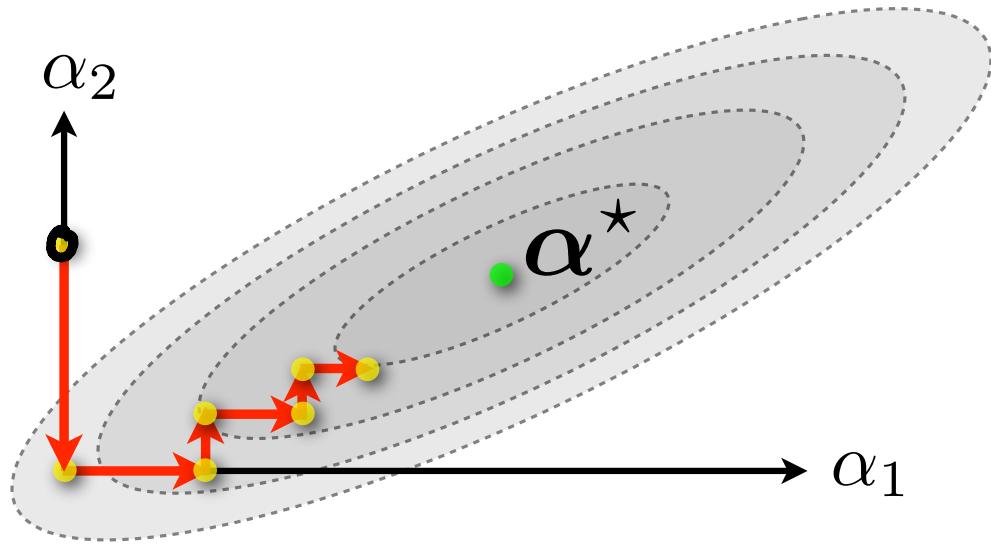


# Coordinate Descent

**Goal:** Find  $\boldsymbol{\alpha}^* \in \mathbb{R}^N$  maximizing or minimizing  $g(\boldsymbol{\alpha})$ .

Yet another optimization algorithm?

**Idea:** Update one coordinate at a time, while keeping others fixed.



initialize  $\boldsymbol{\alpha}^{(0)} \in \mathbb{R}^N$

**for**  $t = 0$ :maxIter **do**

sample

 a coordinate  $n$  randomly from  $1 \dots N$ .

optimize

 $g$  w.r.t. that coordinate:

$$u^* \leftarrow \arg \min_{\textcolor{red}{u} \in \mathbb{R}} \textcolor{blue}{g}(\alpha_1^{(t)}, \dots, \alpha_{n-1}^{(t)}, \textcolor{blue}{u}, \alpha_{n+1}^{(t)}, \dots, \alpha_N^{(t)})$$

update

 $\alpha_n^{(t+1)} \leftarrow u^*$   
 $\alpha_{n'}^{(t+1)} \leftarrow \alpha_{n'}^{(t)}$  for  $n' \neq n$     (unchanged)

**end for**

---

<sup>1</sup>The pseudocode here is for coordinate **descent**, that is to minimize a function. For the equivalent problem of maximizing (coordinate **ascent**), either change this line to arg max, or use the arg min of minus the objective function.

# Issues with SVM

- There is no obvious probabilistic interpretation of SVM.
- Extension to multi-class is non-trivial  
(see Section 14.5.2.4 of KPM book).